

# Evolutionary Computation: Where We Are and Where We're Headed

Kenneth De Jong

Computer Science Department

George Mason University

[kdejong@gmu.edu](mailto:kdejong@gmu.edu)

[www.cs.gmu.edu/~eclab](http://www.cs.gmu.edu/~eclab)

# Historical roots

- Several parallel independent developments in the 1960s.
- Focus: Nature-inspired problem solving
  - i.e., computer science and engineering
- Not modeling of natural systems
  - Interesting area, but not the initial focus

# Common Inspiration

- Problem solving based on a Darwinian notion of an evolutionary system.
- Basic elements:
  - a population of “individuals”
  - a notion of “fitness”
  - a birth/death cycle biased by fitness
  - a notion of “inheritance”

# Historical roots:

- **Evolution Strategies (ESs):**
  - developed by Rechenberg, Schwefel, etc. in 1960s.
  - focus: real-valued parameter optimization
  - individual:     vector of real-valued parameters

# Historical roots:

- **Evolutionary Programming (EP):**
  - Developed by Fogel in 1960s
  - Goal: evolve intelligent behavior
  - Individuals: finite state machines

# Historical roots:

- **Genetic Algorithms (GAs):**
  - developed by Holland in 1960s
  - goal: robust, adaptive systems
  - used an internal “genetic” encoding

# Resulted in:

- wide variety of evolutionary algorithms (EAs)
- wide variety of applications
  - optimization
  - search
  - learning, adaptation
- well-developed analysis
  - theoretical
  - experimental

## **Also resulted in:**


- A bewildering variety of algorithms.
- Historically inconsistent terminology.
- Hard to see relationships, assess strengths & weaknesses, make choices, ...



# A Personal Interest:

- Develop a unifying framework that:
  - Helps one compare and contrast approaches.
  - Encourages crossbreeding.
  - Facilitates intelligent design choices.

# Viewpoint:



# A Simple Evolutionary Algorithm:

1. Randomly generate an initial population.

2. Do until some stopping criteria is met:

Select individuals to be parents (biased by fitness).

Produce offspring.

Select individuals to die (biased by fitness).


End Do.

3. Return a result.


# Simple Example

- Individuals have 2 traits:  $\langle t_1, t_2 \rangle$
- Each trait can vary from -25 to 20
- Unknown fitness surface defined over this trait space.
- Goal: find highly fit individuals quickly


# Generation 0 (random)




# Generation 5



# Generation 10




# Generation 20







# Generation 30




# Generation 0




# Generation 5




# Generation 10



# Generation 20



# Generation 30



# Important Properties:

- Parallel, adaptive search
- Competition for resources:
  - survival (space)
  - reproduce (cpu)
- Stochastic
- Strong sense of fitness “optimization”

# Unified Perspective

- EA as a meta-heuristic to be instantiated:
  - **Core:** a parallel adaptive search heuristic.
  - **Instantiation decisions** to match problem characteristics.
    - Adopt favorable biases



# Key Instantiation Decisions:

- How to represent the space to be searched.
- How to generate “interesting” new samples from existing ones.
- How to choose population sizes.
- How elitist (greedy) should selection be.
- **Answers: my tutorial!**

# Application Areas

- Historically dominant area:  
Parameter-oriented optimization
- Other interesting areas:
  - Searching structure spaces
  - Machine learning
  - Adaptation

# Evolutionary Optimization:

- fitness: function to be optimized
- individuals: points in the space
- reproduction: generating new sample points from existing ones.

# Useful Optimization Properties:

- applicable to continuous, discrete, mixed optimization problems.
- no *a priori* assumptions about convexity, continuity, differentiability, etc.
- relatively insensitive to noise
- easy to parallelize

# Example Areas

- Parameter optimization
  - Mixed parameter data types
  - High dimensional, highly multimodal
- Discrete optimization
  - TSP, SAT, JSS, ...
- Multi-objective optimization
  - Pareto optimality

# Beyond Parameter Optimization

- Search spaces that are not easily parameterized:
  - Data structures
  - Program spaces
  - Conceptual design spaces
  - ...

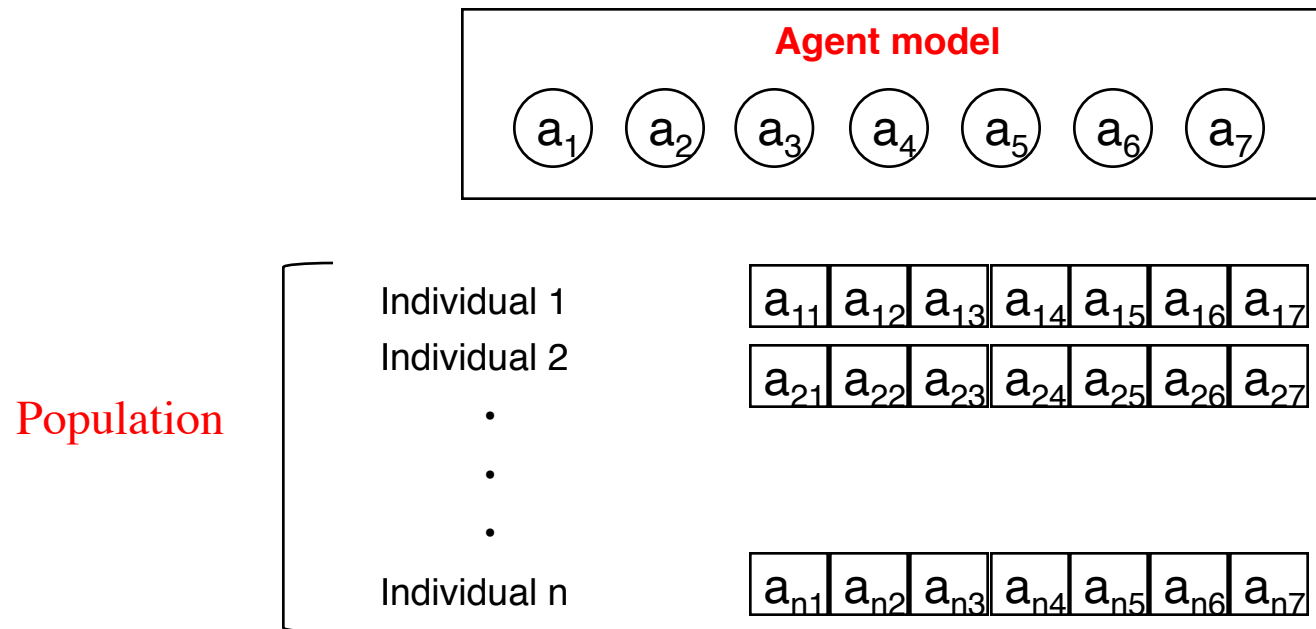
Example:

# Evolving agent behaviors

- What internal elements are evolving?
  - Parameters
  - Structures
  - Code
- What are the mechanisms of change?
  - Mutation
  - Recombination
  - ?

# 1. Evolving agent parameters:

- Identify key parameters that control agent's behavior

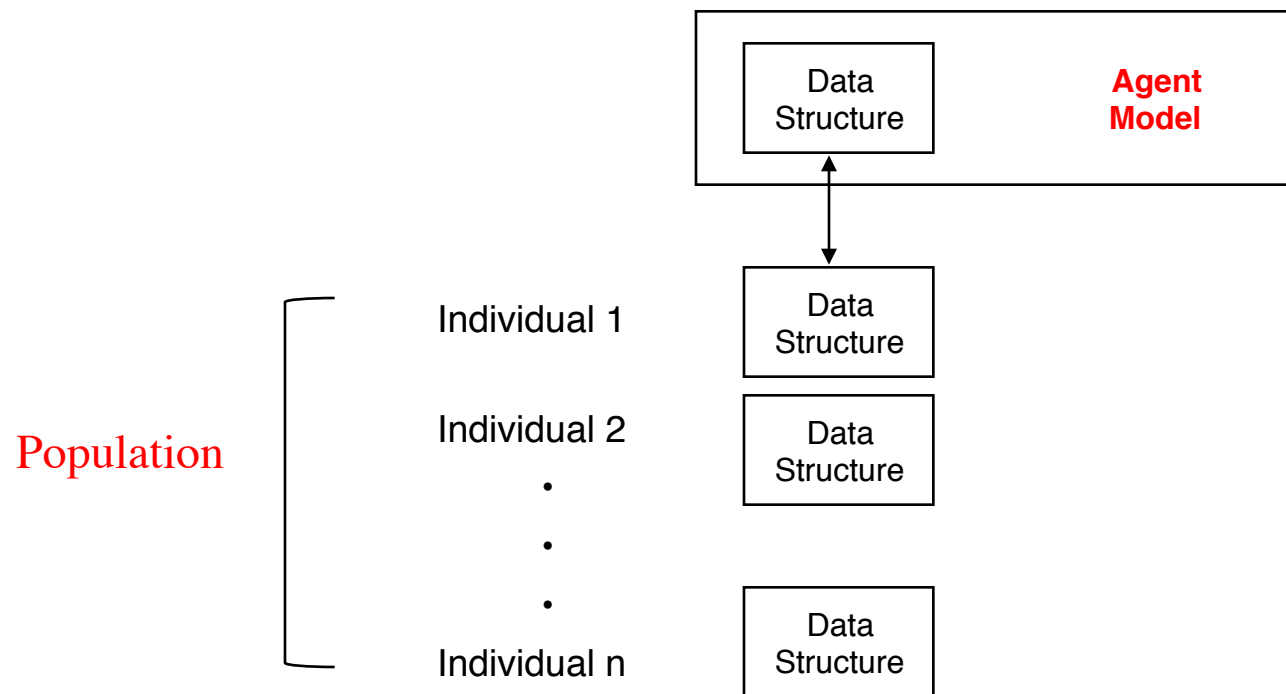


Genome: parameter set  
Mutation: alter individual parameters  
Recombination: combine parameters from different parents



## 2. Evolving agent data structures:

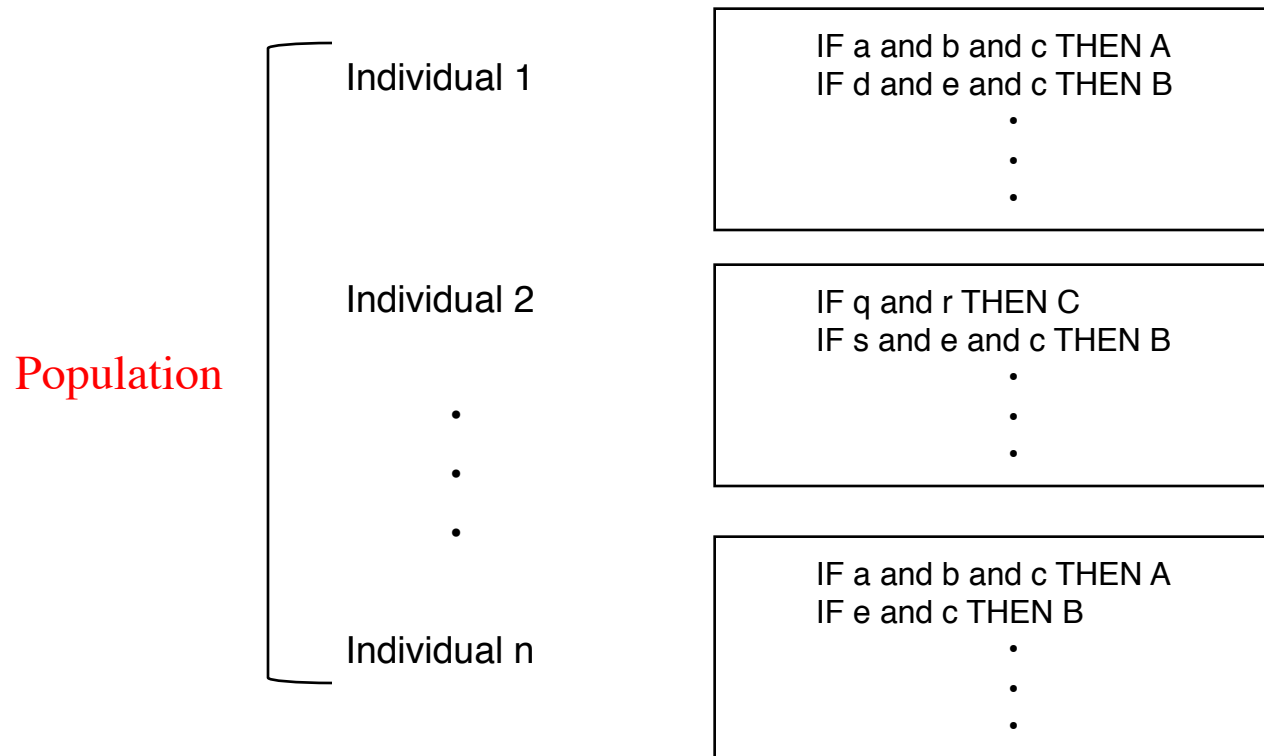
- Modify internal structures that constrain behavior



- Genome represents a graph, tree, ...
- Requires specialized recombination/mutation operators

# 3. Evolving agent programs:

- Modify behavioral procedures



- Genome: entire program
- Recombination: combine rules from different parents
- Mutation: alter individual rules


# Evolving Executable Objects

- Central issue:
  - Evolution-friendly programming languages?
- Syntax, semantics amenable to evolutionary change?
  - Modularity, composability, ...
- Candidates:
  - Lisp, Rules, ANNs, CAs, FSMs, ...

# Example: Evolutionary Robotics

- Collaboration with Naval Research Lab
- Goal: evolve autonomous agent behaviors
- Approach:
  - represent behaviors as rule-based programs
- Behaviors evolved off-line via simulation
- High-fitness behaviors downloaded to robots

# Collision Avoidance and Navigation




- Goal:
  - get a single agent to reliably perform complex navigation tasks.
- Approach:
  - Evolve behaviors offline via simulation
  - Download & test on real robot

# **Navigation and Collision Avoidance**

**Evolved Behaviors**

# Evolving Herding Behavior:



- Goal:
  - Evolve sheep dog herding skills
- Approach:
  - Evolve behaviors offline via simulation
  - Download & test on real robot

# Herding

Evolved Behavior




# Adapting to Partial System Failures

Monitor detects change in internal systems.

Robot performs simulation in virtual world to learn new behaviors.

Successful behaviors performed on-line.






# Evolving Multiple Cooperating Agents

- Example task domains:
  - micro airplane surveillance tasks
  - robo soccer


## Example: Micro Air Vehicles



MAV specs:

- 6-12 inch wingspan
- 50 - 100 grams gross mass

# Goal: evolve collective behaviors for teams of MAVs



# Multi-Agent Evolutionary Design

- RoboCup soccer:
  - [www.cs.gmu.edu/~robotics](http://www.cs.gmu.edu/~robotics)



Skeptical?

Concerned?




Embarrassing moments at gene parties

# Adaptive Testing EAs:

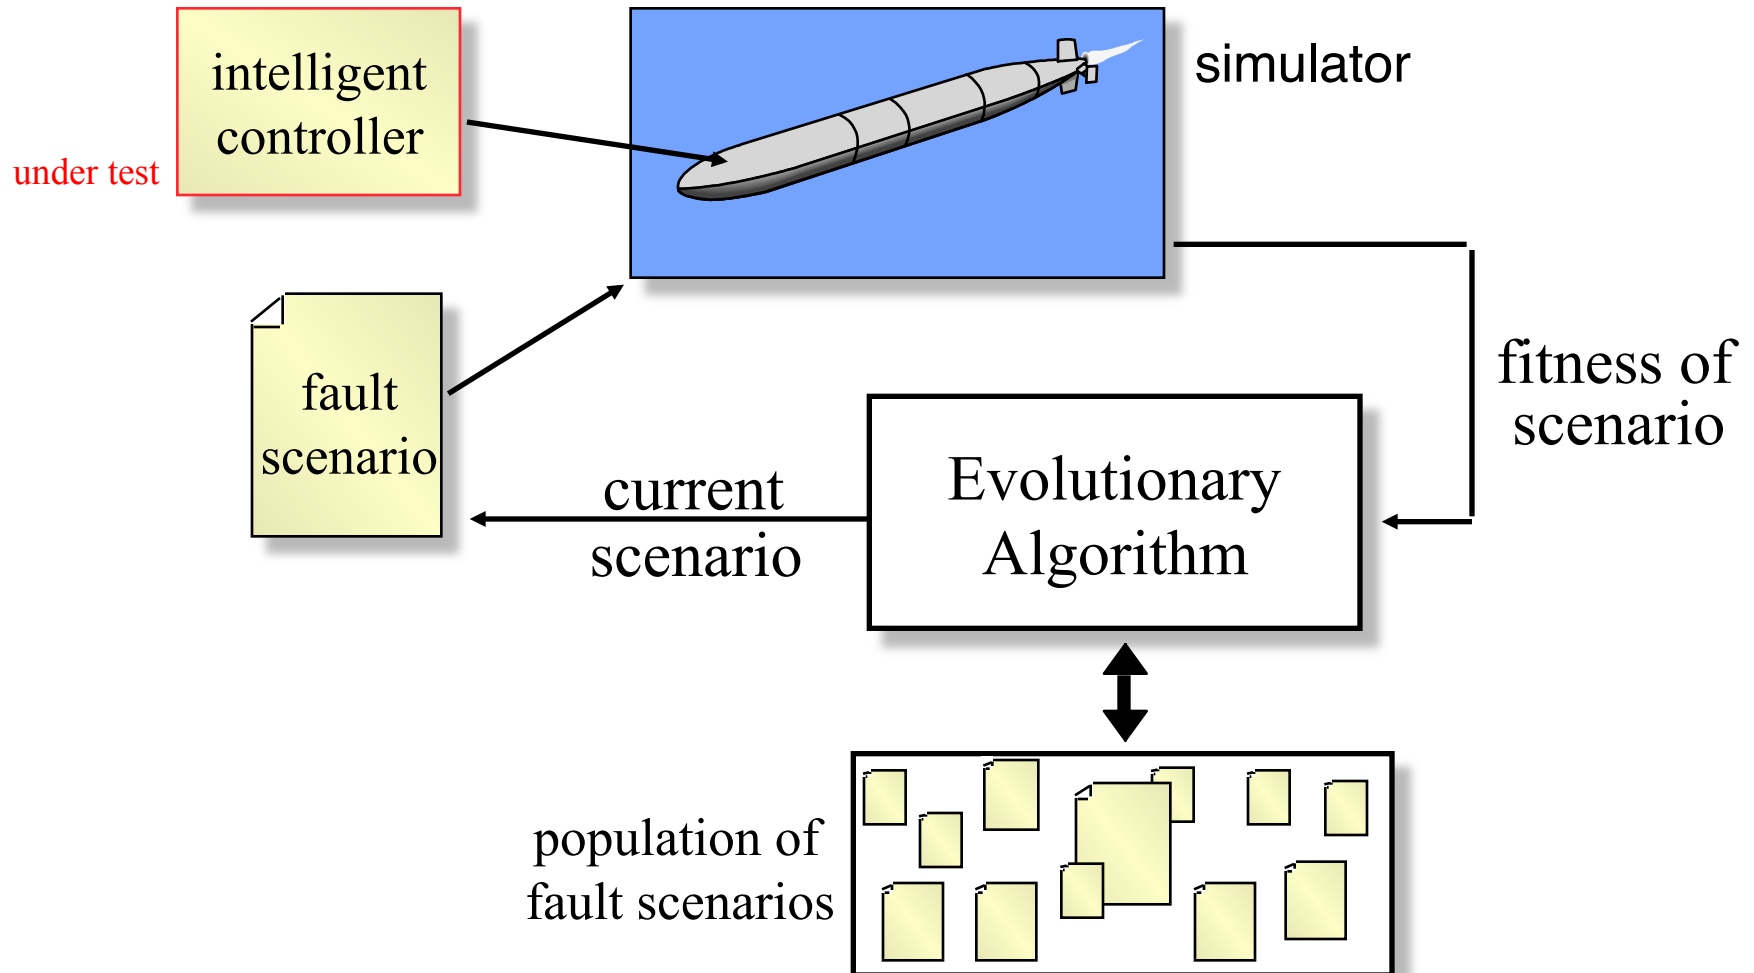
- How to validate autonomous systems?
  - Prove theorems?
  - Hire test engineers?
- Interesting alternative:
  - Use EAs to evolve new scenarios.
  - Scenario's fitness related to the difficulties it creates for autonomous agents.




# The Traditional Test Cycle



# Adaptive Testing



# Evolutionary Testing of Draper Labs AUV



- Very high fidelity simulation of autonomous underwater vehicle.
- Modeled real vehicle (Draper/Darpa).


# Evolutionary Testing of AUV:

- Controller failures found included:

- Vehicle exceeds maximum roll rate, oscillates, etc.
- Vehicle crashes into bottom
- Vehicle unable to complete mission in time

- Some missions demonstrated failures that had not been observed before.

- Fault scenarios judged very interesting by controller and vehicle design teams.



# Broader context:

- How to understand/test/evaluate complex systems?
- Interesting approach: **ABM+EC**
  - Build an **agent-based model** of a complex system
  - Use **EAs** to evaluate, test, etc.
- Developed tools:
  - **ECJ**: EC toolkit
  - **MASON**: multi-agent simulation toolkit

# Examples:

- Network security (evolve hackers)
- Homeland security (evolve terrorist scenarios)
  - Water supply
  - Power grids
  - ...

# Terrorist Threats to Existing Water Systems

- Biological or chemical contamination is placed in a water system.
- Large number of system's nodes are affected.
- Goal: to maximize the negative impact measured by the number of nodes affected.

# Evolutionary Testing

- **ABM:** Combo of MASON model and EPANet models of existing water systems.
- **ECJ:** Evolved terrorist scenarios that maximized simulated damage.




# Initial Findings

- Identified several serious weaknesses
- Gave focus for limited budget remedies
- Useful for what-if scenarios

# Computer Network Security

- Starting point: a AB network intrusion model
  - 2500 potentially vulnerable systems
  - Modeled Blue force (security) responses
  - Modeled Red force (hacker) tools
- Evolved Red force tactics over time

# Evolving Hacker Performance



# “In Silico” Science

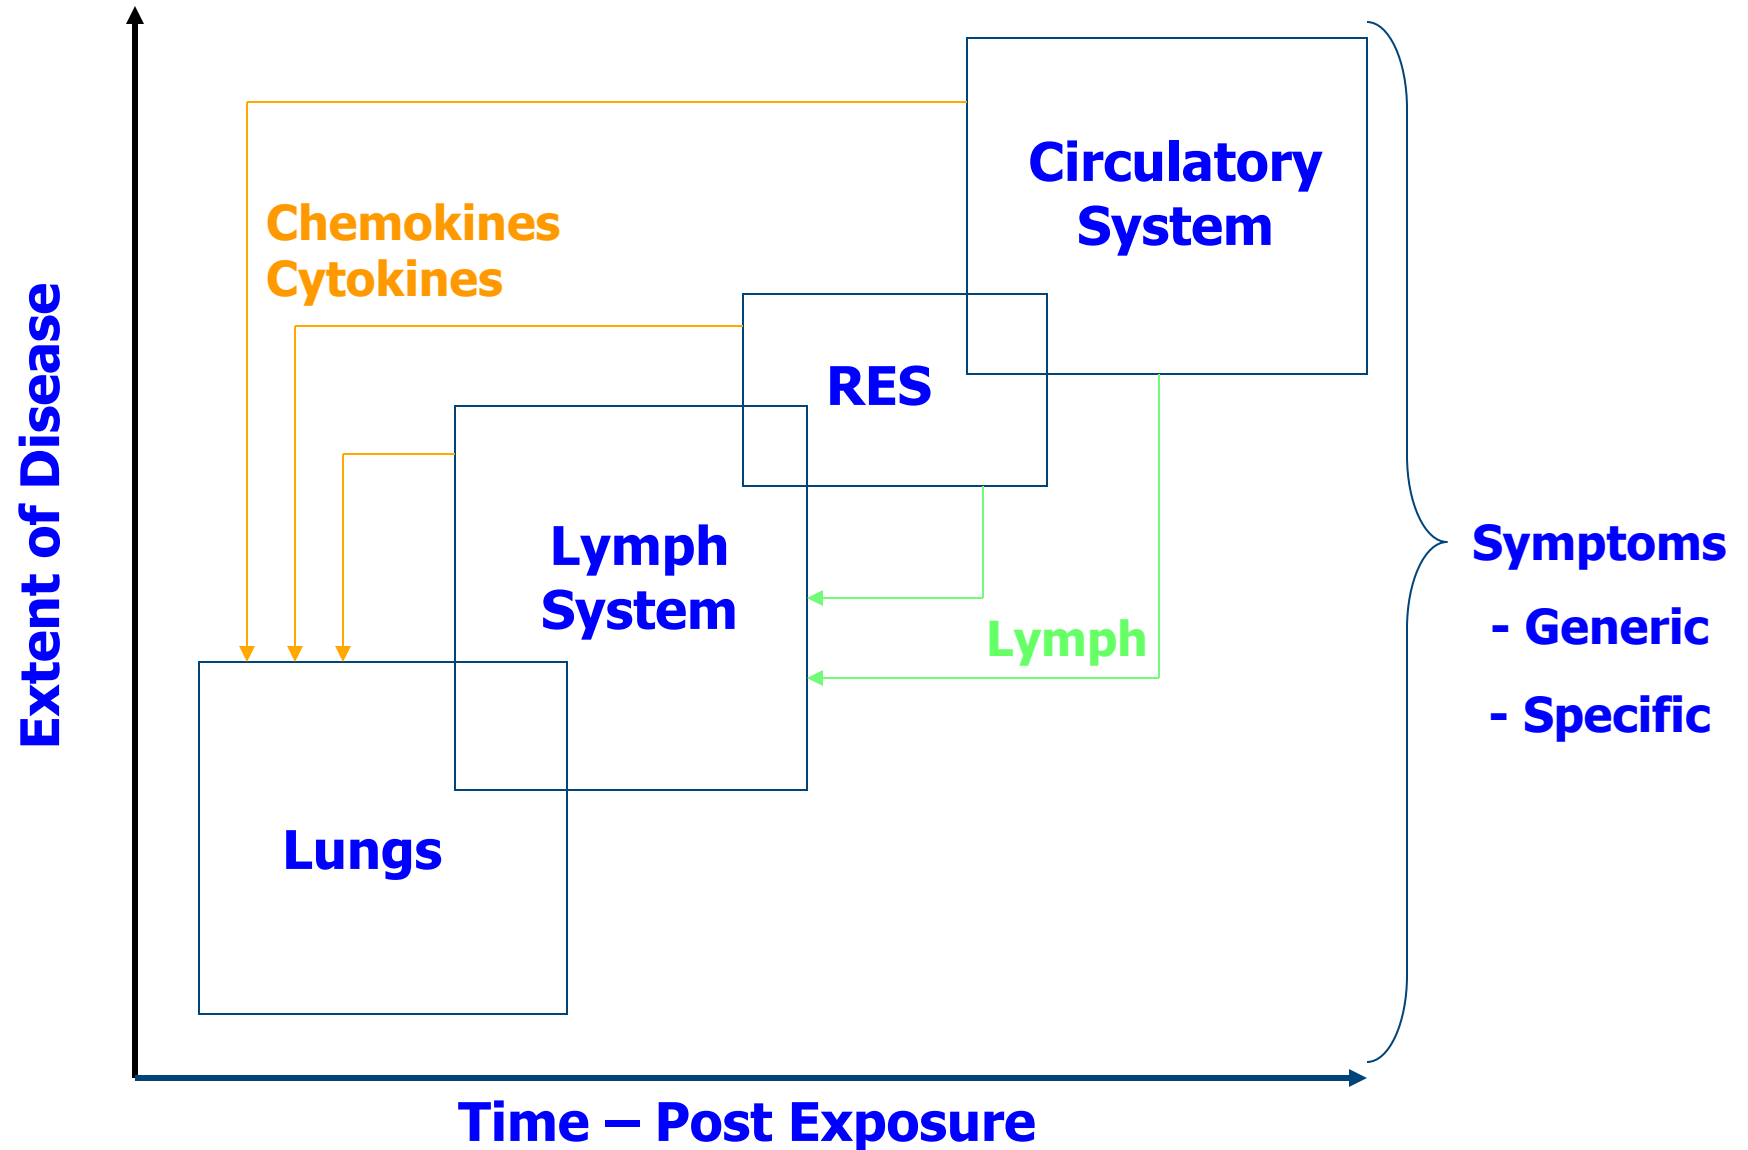
- 3<sup>rd</sup> way of doing science:
  - *In vitro*, *in vivo*, ...
- Serious alternative to:
  - Experiments with animals
  - Experiments with humans

Example:

# Understanding Inhalation Anthrax

- Very real concerns: 1990 exposures
- Clear lack of understanding
- Clear opportunity for *in silico* approaches

# Developed an AB Anthrax Model



# Applying EAs

- Explore space of treatment scenarios
  - Suggest new interventions
- Hypothesis generation:
  - Suggest new lab experiments

# Summary: Unified framework

- Meta-heuristic perspective:
  - Consider properties of the objects to be evolved
  - Select internal representation
  - Choose reproductive operators
  - Choose population sizes
  - Choose selection pressure
- Result: a well-designed EA



# However ...

- New applications pressing state of the art.
- Unified view of “simple EAs” is not sufficient.
- Principled extensions are required.

# Additional inspiration from nature?

- Relationship of simple EAs to Biology
  - Fairly remote!
- E.g.,
  - No notion age, growth, development
  - No distinction between genotype and phenotype
  - No notion of gender, multiple species
  - ...

# Relationship of simple EAs to Biology

- Clear tension between:  
Biological fidelity and computational efficiency
- New directions:
  - Exploring other biological aspects that improve computational efficiency.

# EC direction: **Adaptive EAs**

Goal: reduce tuning efforts

- EAs have their own parameters
- How to tune them?
  - Manually?
  - Meta-EAs?
- Far better to have:
  - **Self-adapting mechanisms**

# EC direction: Adaptive EAs

## Adaptive strategies:

- Adapt across multiple runs/restarts
- Adapt during a single run:
  - Via feedback control mechanisms
    - E.g., adaptive Gaussian mutation
  - Via evolutionary mechanisms
    - E.g., parameters part of the genome

# EA direction: Exploiting Parallelism

Low-hanging fruit: parallel evaluation

- E.g., master-slave architectures
- Tougher challenges:
  - Coarsely-grained distributed systems
    - E.g., across beowulf clusters
  - Finely-grained multi-threaded systems
    - E.g., exploit gpu clusters

# Coarsely-grained parallelism

- EA Island models:
  - Multiple EAs running on separate machines
  - Heterogeneous EAs?
  - Occasional migrations to other machines
- Challenges:
  - Clearly different than single population models
  - When to use them?
  - How to configure them?

# Fine-grained parallelism


- EA cellular models:
  - Introduce a spatial topology in the population
  - Micro EA in each cell
  - Only local interactions with neighbors
  - All micro EAs run in parallel
- Challenges:
  - Clearly different than single population models
  - When to use them?
  - How to configure them?



EC direction: Multi-objective optim.

- Impressive initial results
- Need to scale up to  $>3$ -4 objectives
- Role of more complex EAs
- Need deeper theoretical understanding


# EC direction: **Non-stationarity**



- Time-varying environments:
  - fitness landscape changes during evolution
  - goal: adaptation, tracking
  - standard optimization-oriented EAs not well-suited for this.

# EC direction: Multiple Species

- Speciation: Non-random mating
  - Helps maintain diversity
  - Simultaneously explore multiple peaks
  - Issues:
    - How, when, how many?
    - Emergent?



# EC direction: Co-evolution

- Multiple populations/species
- Fitness is a function of other populations/species
- Examples:
  - Competitive co-evolution
  - Cooperative co-evolution

# Competitive co-evolution

- Improve performance via “arms race”
  - E.g., Hillis’ sorting networks
  - E.g., competitive games
- Challenges:
  - When to use?
  - How to configure?

# Cooperative co-evolution

- Improve performance via “teamwork”
  - Decompose problem into subcomponents
  - Evolve subcomponents in parallel
  - Fitness is a function of other components
- Challenges:
  - When to use?
  - How to configure?

# EC direction: Morphogenesis

- Inspiration from biology:
  - Strong distinction between:
    - Genotype (plans)
    - Phenotype (objects)
- Evolve plans, not objects
- Morphogenesis: Generate objects from plans

# Example: Evolutionary Design

- Standard EAs are good for parameter optimization of designs.
- Different EAs are needed for conceptual design.



# Phenotypic Representation

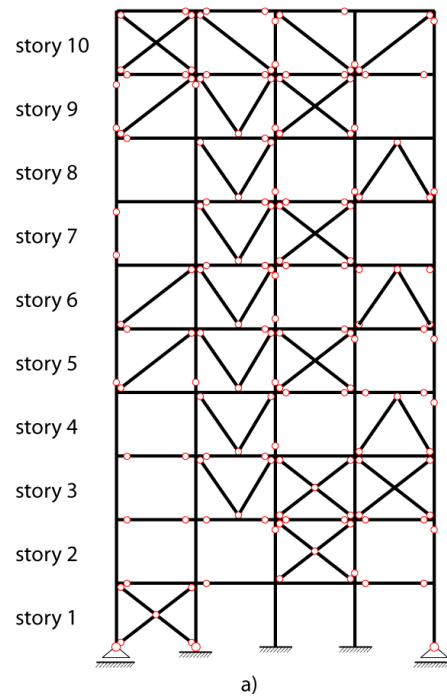
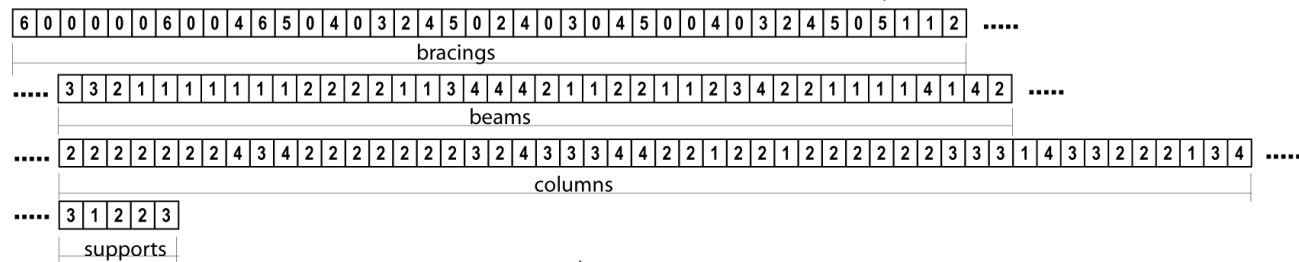




Diagram (b) shows a 10x10 matrix representing the phenotypic representation of the building frame. The matrix is symmetric and contains numerical values representing the properties of the frame elements.

		-4		-1		-4		-2	
2	5	2	1	1	1	3	2	4	
	-1		-1		-1		-1		
1	2	4	4	3	5	3	0	2	
	-3		-4		-2		-2		
2	0	2	4	3	0	3	3	3	
	-2		-1		-1		-2		
1	0	2	4	2	5	2	0	2	
	-2		-1		-1		-2		
2	2	2	4	1	0	2	3	2	
	-3		-4		-4		-4		
3	2	3	4	3	5	4	0	4	
	-2		-2		-1		-1		
2	0	2	4	3	0	2	3	4	
	-1		-1		-2		-2		
2	0	2	4	2	6	2	5	2	
	-1		-1		-1		-1		
2	0	2	0	4	6	3	0	4	
	-3		-3		-2		-1		
2	6	2	0	2	0	2	0	2	
3		1		2		2		3	



c)

# Generative Representation



# Generative Representations



- “Inspiration from nature” strategy has yielded a variety of generative models:
  - L-systems
  - Cellular automata
  - Genetic regulatory networks
  - ...

# L-systems

- Short for Lindenmayer systems
  - Developed ~ 1968 by A. Lindenmayer
  - Goal: model plant morphology
  - Approach: formal “rewrite” grammars:


# L-systems

- Examples:



# L-systems

- Evolutionary design:
  - Evolve set of rewrite rules to achieve a goal.
- Examples:
  - Artificial vegetation (Ochoa, Jacob, Moch, ...)
  - Tables (Popovici)
  - Robot morphology (Hornby, Pollack, Bentley, ...)
  - Neural networks (Kitano, ...)



# Cellular Automata



- Similar in spirit to L-systems:
  - Symbols, axioms, rewrite rules
  - Explicit bounded spatial topology
- Classical examples:
  - Conway's Game of Life
  - Wolfram's pattern generators

# Cellular Automata

- Elements:
  - A topology of cells
  - Initial state of each cell
  - State transition rewrite rules
    - Based on neighboring states
- Simple rules generate emergent complexity



# Cellular Automata



# Cellular Automata

- Evolutionary design:

Evolve transition rules to achieve a goal

- Examples:

- Artificial societies (Axtell, Ilichinski, ...)
- Programs (Mitchell, Crutchfield, Sipper, ...)
- Chemical structures (Gerhardt, Schuster, ...)
- Building designs (Kicingier, Arciszewski, ...)

# Genetic Regulatory Networks (GRNs)

- Similar in spirit to L-systems, CAs:
  - Symbols, axioms, production rules
  - Explicit gene/cell model
    - Genes: independent rewrite rules
    - Cells: collections of genes + shared global memory
      - Similar to AI “blackboard” models

# GRNs

- Evolutionary design:


Evolve production rules to achieve a goal

- Examples:

- Robot control (Kumar, Grajdeanu, ...)
- Multi-cellular objects (Miller, Federici, Gordon, ...)

# GRNs

- Example: obstacle avoidance



# GRNs

- Example: artificial embryogeny
  - Seed: single cell
  - Morphogenesis (via embryogeny):
    - Model cell division, growth, maturation
  - Grow complex, multidimensional objects



# Generative Representations

- Interesting open EC question:
  - Which generative representations are more evolution-friendly?
    - L-systems
    - Cellular automata
    - Genetic regulatory networks
    - ...

# EC direction: **Agent orientation**

- Individuals more autonomous, active
- Fitness is a function of other agents and environment-altering actions
- E.g.,
  - Evolutionary robotics
  - Evolution of cooperation



# EC direction: **Analysis**

- Need stronger analysis tools:
  - Markov models
  - Statistical mechanics
  - Evolutionary game theory
  - Test problem generators
  - Visualization

# EC direction: Hybrid systems

- Continuing to explore:
  - Memetic algorithms: EAs and local search
  - COGANNs: EAs and ANNs
  - EAs and symbolic machine learning
  - EAs and agent-based models
  - ...

# EA Generalizations:

- Meta-heuristics:
  - Heuristics for designing heuristics
    - E.g., hill climbing, greedy, ...
  - Adopt no-free lunch view
  - Instantiate templates in a problem-specific manner

# EA Generalizations:

- Nature-Inspired Computation:
  - Early example: simulated annealing
  - Today: evolutionary algorithms
  - Others:
    - particle swarm optimization
    - ant colony optimization
    - ...

# Conclusions:

- We've come a long way.
- Many new challenges ahead.
- A strategy for continued success:
  - An expanded unified framework that leads to:
    - Principled design
    - Principled extensions

# More information:

- Journals:
  - Evolutionary Computation (MIT Press)
  - Trans. on Evolutionary Computation (IEEE)
  - Genetic Programming & Evolvable Hardware
- Conferences:
  - GECCO, CEC, PPSN, FOGA, ...
- Internet:
  - [www.cs.gmu.edu/~eclab](http://www.cs.gmu.edu/~eclab)
- Lots of books including:
  - Evolutionary Computation: A Unified Approach
    - Kenneth De Jong, MIT Press, 2006

